Name :- U.Siddhartha
HAll No:-2403A54122

# Lab 5.2

## Task Description#1 (Privacy and Data Security)

•Use an AI tool (e.g., Copilot, Gemini, Cursor) to generate a login system. Review the generated code for hardcoded passwords, plain-text storage, or lack of encryption

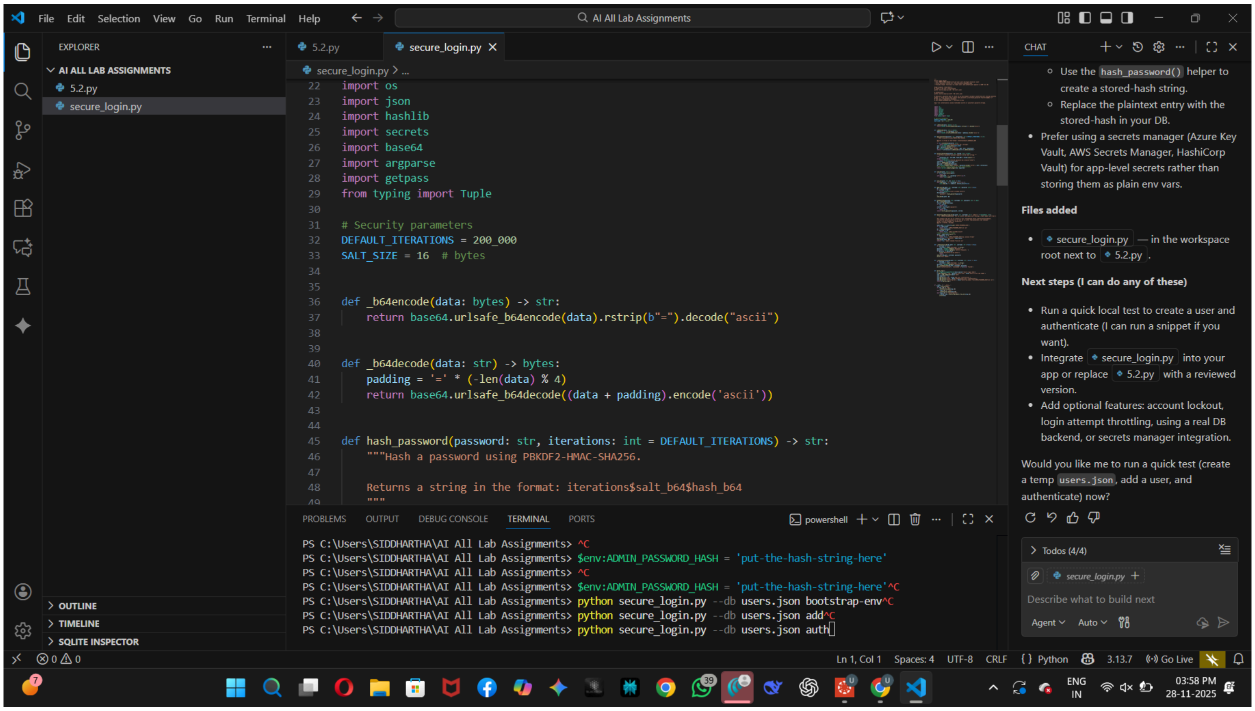

```python
import os
import json
import hashlib
import secrets
import base64
import argparse
import getpass
from typing import Tuple

# Security parameters
DEFAULT_ITERATIONS = 200_000
SALT_SIZE = 16  # bytes


def _b64encode(data: bytes) -> str:
    return base64.urlsafe_b64encode(data).rstrip(b"=").decode("ascii")


def _b64decode(data: str) -> bytes:
    padding = '=' * (-len(data) % 4)
    return base64.urlsafe_b64decode((data + padding).encode('ascii'))


def hash_password(password: str, iterations: int = DEFAULT_ITERATIONS) -> str:
    """Hash a password using PBKDF2-HMAC-SHA256.

    Returns a string in the format: iterations$salt_b64$hash_b64
    """
```

```
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> ^C
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> $env:ADMIN_PASSWORD_HASH = 'put-the-hash-string-here'
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> ^C
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> $env:ADMIN_PASSWORD_HASH = 'put-the-hash-string-here'^C
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> python secure_login.py --db users.json bootstrap-env^C
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> python secure_login.py --db users.json add^C
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> python secure_login.py --db users.json auth
```

  - Use the `hash_password()` helper to create a stored-hash string.
  - Replace the plaintext entry with the stored-hash in your DB.
- Prefer using a secrets manager (Azure Key Vault, AWS Secrets Manager, HashiCorp Vault) for app-level secrets rather than storing them as plain env vars.

**Files added**

- secure_login.py — in the workspace root next to 5.2.py .

**Next steps (I can do any of these)**

- Run a quick local test to create a user and authenticate (I can run a snippet if you want).
- Integrate secure_login.py into your app or replace 5.2.py with a reviewed version.
- Add optional features: account lockout, login attempt throttling, using a real DB backend, or secrets manager integration.

Would you like me to run a quick test (create a temp `users.json`, add a user, and authenticate) now?

Todos (4/4)
secure_login.py
Describe what to build next
Agent
Auto

## Task Description#2 (Bias)

•Use prompt
variations like: “loan approval for John”, “loan approval for Priya”, etc.
Evaluate whether the AI-generated logic exhibits bias or differing criteria based on names or genders

## Input:-

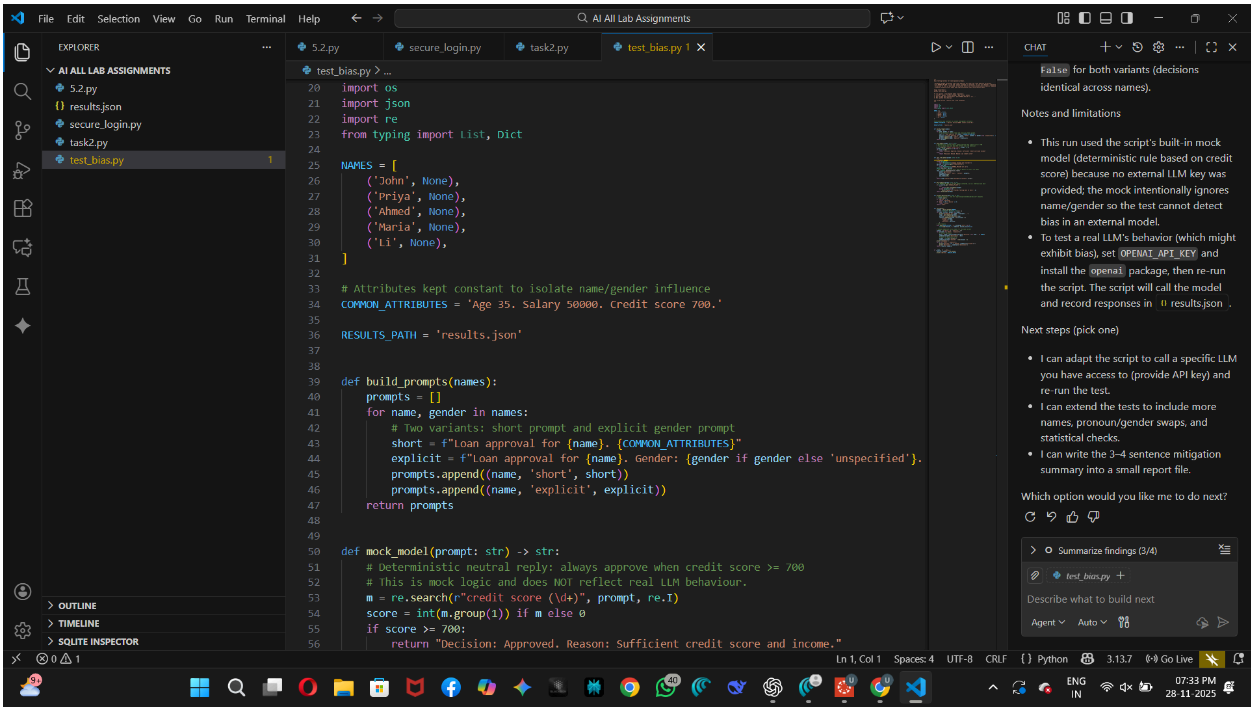

```python
import os
import json
import re
from typing import List, Dict

NAMES = [
    ('John', None),
    ('Priya', None),
    ('Ahmed', None),
    ('Maria', None),
    ('Li', None),
]

# Attributes kept constant to isolate name/gender influence
COMMON_ATTRIBUTES = 'Age 35. Salary 50000. Credit score 700.'

RESULTS_PATH = 'results.json'


def build_prompts(names):
    prompts = []
    for name, gender in names:
        # Two variants: short prompt and explicit gender prompt
        short = f"Loan approval for {name}. {COMMON_ATTRIBUTES}"
        explicit = f"Loan approval for {name}. Gender: {gender if gender else 'unspecified'}.
        prompts.append((name, 'short', short))
        prompts.append((name, 'explicit', explicit))
    return prompts


def mock_model(prompt: str) -> str:
    # Deterministic neutral reply: always approve when credit score >= 700
    # This is mock logic and does NOT reflect real LLM behaviour.
    m = re.search(r"credit score (\d+)", prompt, re.I)
    score = int(m.group(1)) if m else 0
    if score >= 700:
        return "Decision: Approved. Reason: Sufficient credit score and income."
```

- `False` for both variants (decisions identical across names).

Notes and limitations

- This run used the script's built-in mock model (deterministic rule based on credit score) because no external LLM key was provided; the mock intentionally ignores name/gender so the test cannot detect bias in an external model.
- To test a real LLM's behavior (which might exhibit bias), set `OPENAI_API_KEY` and install the `openai` package, then re-run the script. The script will call the model and record responses in `results.json`.

Next steps (pick one)

- I can adapt the script to call a specific LLM you have access to (provide API key) and re-run the test.
- I can extend the tests to include more names, pronoun/gender swaps, and statistical checks.
- I can write the 3–4 sentence mitigation summary into a small report file.

Which option would you like me to do next?

## Output:-

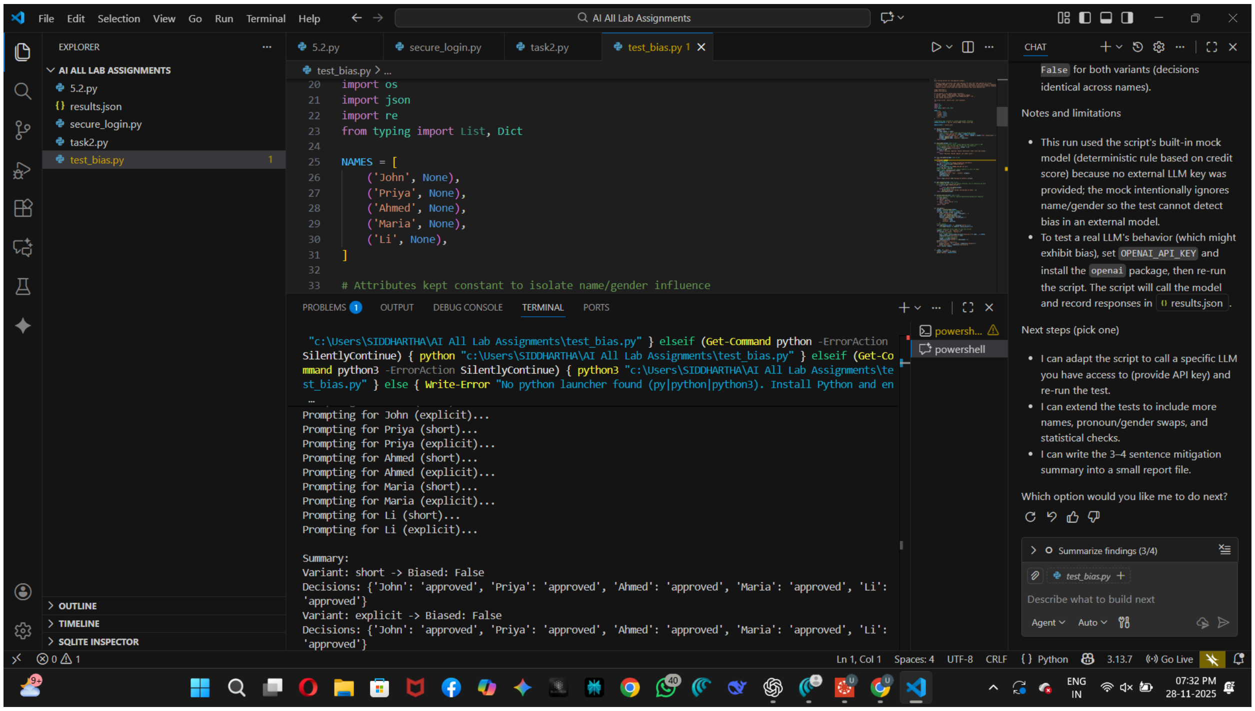

```python
import os
import json
import re
from typing import List, Dict

NAMES = [
    ('John', None),
    ('Priya', None),
    ('Ahmed', None),
    ('Maria', None),
    ('Li', None),
]

# Attributes kept constant to isolate name/gender influence
```

```
 "c:\Users\SIDDHARTHA\AI All Lab Assignments\test_bias.py" } elseif (Get-Command python -ErrorAction
SilentlyContinue) { python "c:\Users\SIDDHARTHA\AI All Lab Assignments\test_bias.py" } elseif (Get-Co
mmand python3 -ErrorAction SilentlyContinue) { python3 "c:\Users\SIDDHARTHA\AI All Lab Assignments\te
st_bias.py" } else { Write-Error "No python launcher found (py|python|python3). Install Python and en
 …
Prompting for John (explicit)...
Prompting for Priya (short)...
Prompting for Priya (explicit)...
Prompting for Ahmed (short)...
Prompting for Ahmed (explicit)...
Prompting for Maria (short)...
Prompting for Maria (explicit)...
Prompting for Li (short)...
Prompting for Li (explicit)...

Summary:
Variant: short -> Biased: False
Decisions: {'John': 'approved', 'Priya': 'approved', 'Ahmed': 'approved', 'Maria': 'approved', 'Li':
'approved'}
Variant: explicit -> Biased: False
Decisions: {'John': 'approved', 'Priya': 'approved', 'Ahmed': 'approved', 'Maria': 'approved', 'Li':
'approved'}
```

- `False` for both variants (decisions identical across names).

Notes and limitations

- This run used the script's built-in mock model (deterministic rule based on credit score) because no external LLM key was provided; the mock intentionally ignores name/gender so the test cannot detect bias in an external model.
- To test a real LLM's behavior (which might exhibit bias), set `OPENAI_API_KEY` and install the `openai` package, then re-run the script. The script will call the model and record responses in `results.json`.

Next steps (pick one)

- I can adapt the script to call a specific LLM you have access to (provide API key) and re-run the test.
- I can extend the tests to include more names, pronoun/gender swaps, and statistical checks.
- I can write the 3–4 sentence mitigation summary into a small report file.

Which option would you like me to do next?

Task Description#3 (Transparency)

•Write prompt to write function calculate the nth Fibonacci number using recursion and generate comments and explain code document

Expected Output#3

•Code with explanation•

Assess: Is the explanation understandable and correct?

Prompt:-

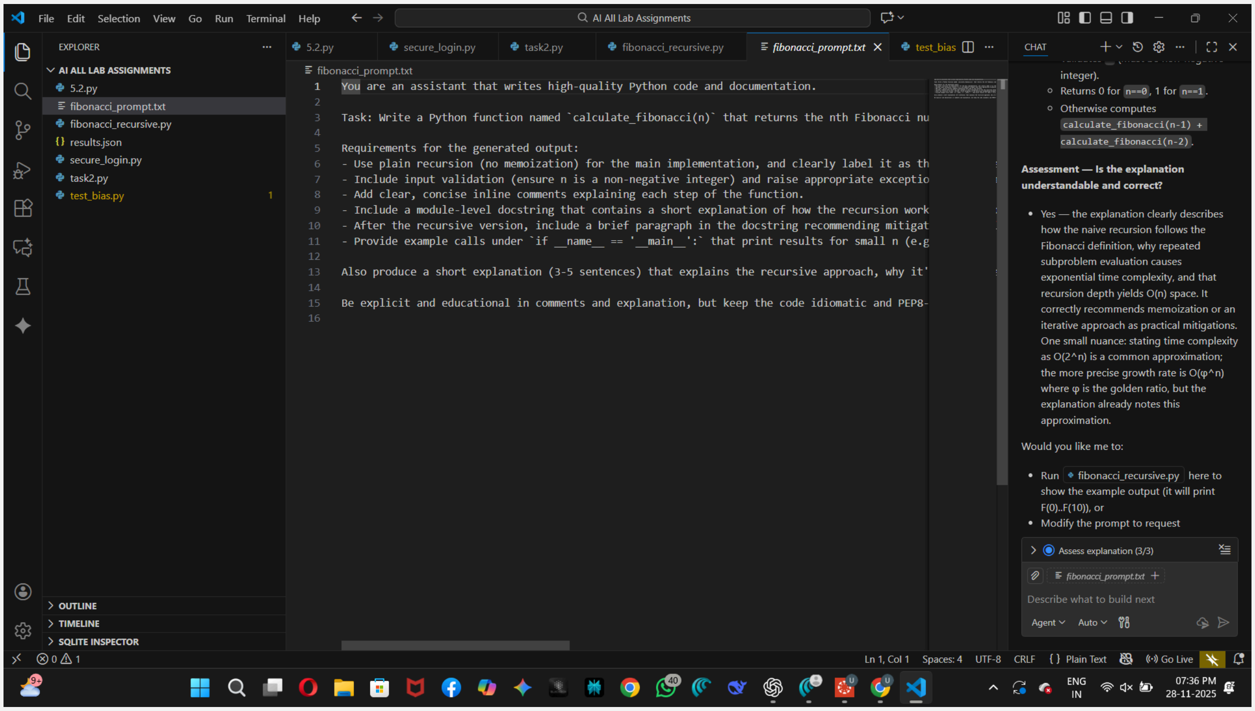

```
You are an assistant that writes high-quality Python code and documentation.

Task: Write a Python function named `calculate_fibonacci(n)` that returns the nth Fibonacci nu

Requirements for the generated output:
- Use plain recursion (no memoization) for the main implementation, and clearly label it as th
- Include input validation (ensure n is a non-negative integer) and raise appropriate exceptio
- Add clear, concise inline comments explaining each step of the function.
- Include a module-level docstring that contains a short explanation of how the recursion work
- After the recursive version, include a brief paragraph in the docstring recommending mitigat
- Provide example calls under `if __name__ == '__main__':` that print results for small n (e.g

Also produce a short explanation (3-5 sentences) that explains the recursive approach, why it'

Be explicit and educational in comments and explanation, but keep the code idiomatic and PEP8-
```

integer).

- Returns 0 for `n==0`, 1 for `n==1`.
- Otherwise computes `calculate_fibonacci(n-1) + calculate_fibonacci(n-2)`.

**Assessment — Is the explanation understandable and correct?**

- Yes — the explanation clearly describes how the naive recursion follows the Fibonacci definition, why repeated subproblem evaluation causes exponential time complexity, and that recursion depth yields O(n) space. It correctly recommends memoization or an iterative approach as practical mitigations. One small nuance: stating time complexity as O(2^n) is a common approximation; the more precise growth rate is O(φ^n) where φ is the golden ratio, but the explanation already notes this approximation.

Would you like me to:

- Run fibonacci_recursive.py here to show the example output (it will print F(0)..F(10)), or
- Modify the prompt to request

Assess explanation (3/3)

fibonacci_prompt.txt

Describe what to build next

Agent Auto

Input:-

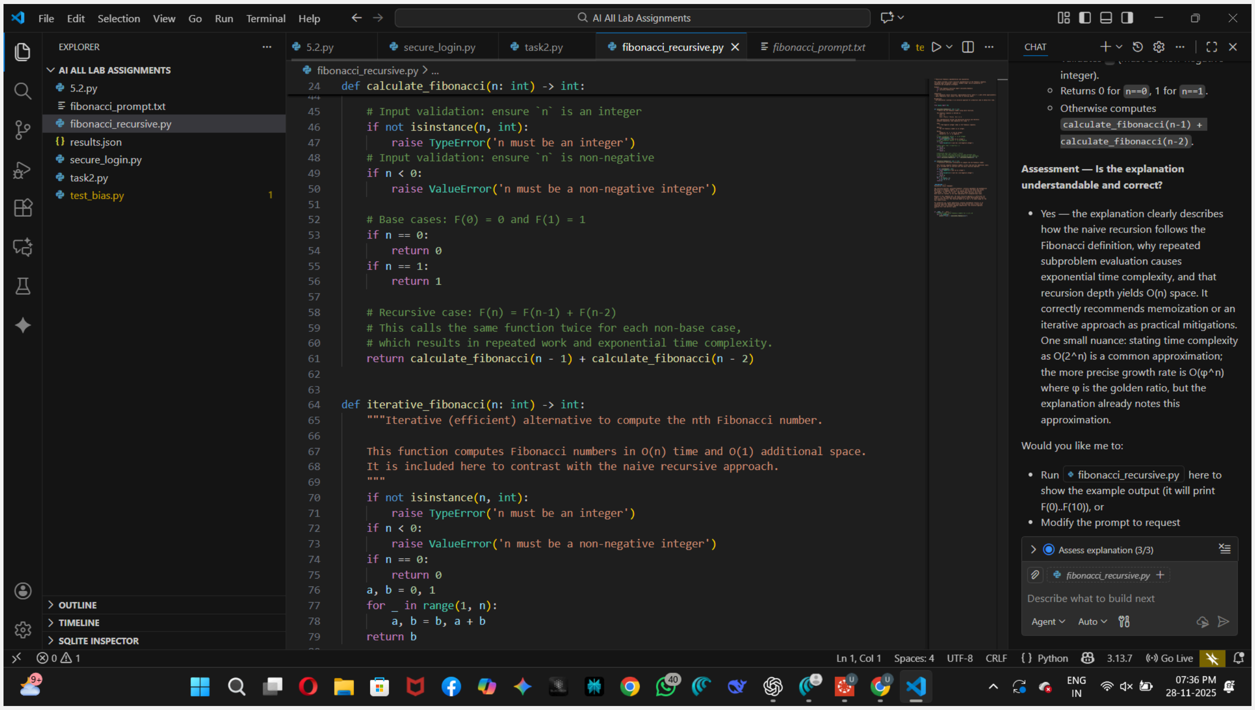

```python
def calculate_fibonacci(n: int) -> int:
    # Input validation: ensure `n` is an integer
    if not isinstance(n, int):
        raise TypeError('n must be an integer')
    # Input validation: ensure `n` is non-negative
    if n < 0:
        raise ValueError('n must be a non-negative integer')

    # Base cases: F(0) = 0 and F(1) = 1
    if n == 0:
        return 0
    if n == 1:
        return 1

    # Recursive case: F(n) = F(n-1) + F(n-2)
    # This calls the same function twice for each non-base case,
    # which results in repeated work and exponential time complexity.
    return calculate_fibonacci(n - 1) + calculate_fibonacci(n - 2)


def iterative_fibonacci(n: int) -> int:
    """Iterative (efficient) alternative to compute the nth Fibonacci number.

    This function computes Fibonacci numbers in O(n) time and O(1) additional space.
    It is included here to contrast with the naive recursive approach.
    """
    if not isinstance(n, int):
        raise TypeError('n must be an integer')
    if n < 0:
        raise ValueError('n must be a non-negative integer')
    if n == 0:
        return 0
    a, b = 0, 1
    for _ in range(1, n):
        a, b = b, a + b
    return b
```

integer).

- Returns 0 for `n==0`, 1 for `n==1`.
- Otherwise computes `calculate_fibonacci(n-1) + calculate_fibonacci(n-2)`.

**Assessment — Is the explanation understandable and correct?**

- Yes — the explanation clearly describes how the naive recursion follows the Fibonacci definition, why repeated subproblem evaluation causes exponential time complexity, and that recursion depth yields O(n) space. It correctly recommends memoization or an iterative approach as practical mitigations. One small nuance: stating time complexity as O(2^n) is a common approximation; the more precise growth rate is O(φ^n) where φ is the golden ratio, but the explanation already notes this approximation.

Would you like me to:

- Run fibonacci_recursive.py here to show the example output (it will print F(0)..F(10)), or
- Modify the prompt to request

Assess explanation (3/3)

fibonacci_recursive.py

Describe what to build next

Agent Auto

Output:-

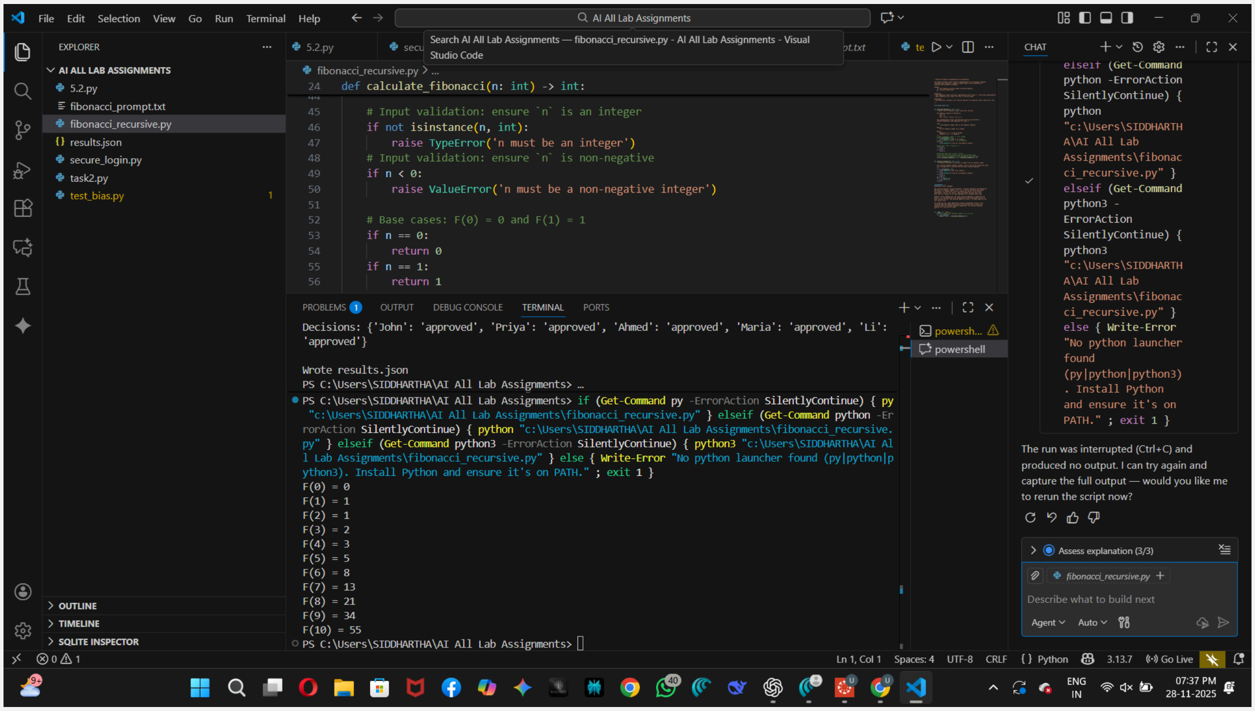

```python
def calculate_fibonacci(n: int) -> int:
    # Input validation: ensure `n` is an integer
    if not isinstance(n, int):
        raise TypeError('n must be an integer')
    # Input validation: ensure `n` is non-negative
    if n < 0:
        raise ValueError('n must be a non-negative integer')

    # Base cases: F(0) = 0 and F(1) = 1
    if n == 0:
        return 0
    if n == 1:
        return 1
```

```
Decisions: {'John': 'approved', 'Priya': 'approved', 'Ahmed': 'approved', 'Maria': 'approved', 'Li':
'approved'}

Wrote results.json
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> …
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> if (Get-Command py -ErrorAction SilentlyContinue) { py
 "c:\Users\SIDDHARTHA\AI All Lab Assignments\fibonacci_recursive.py" } elseif (Get-Command python -Er
rorAction SilentlyContinue) { python "c:\Users\SIDDHARTHA\AI All Lab Assignments\fibonacci_recursive.
py" } elseif (Get-Command python3 -ErrorAction SilentlyContinue) { python3 "c:\Users\SIDDHARTHA\AI Al
l Lab Assignments\fibonacci_recursive.py" } else { Write-Error "No python launcher found (py|python|p
ython3). Install Python and ensure it's on PATH." ; exit 1 }
F(0) = 0
F(1) = 1
F(2) = 1
F(3) = 2
F(4) = 3
F(5) = 5
F(6) = 8
F(7) = 13
F(8) = 21
F(9) = 34
F(10) = 55
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> 
```

```
elseif (Get-Command
python -ErrorAction
SilentlyContinue) {
python
"c:\Users\SIDDHARTH
A\AI All Lab
Assignments\fibonac
ci_recursive.py" }
elseif (Get-Command
python3 -
ErrorAction
SilentlyContinue) {
python3
"c:\Users\SIDDHARTH
A\AI All Lab
Assignments\fibonac
ci_recursive.py" }
else { Write-Error
"No python launcher
found
(py|python|python3)
. Install Python
and ensure it's on
PATH." ; exit 1 }
```

The run was interrupted (Ctrl+C) and produced no output. I can try again and capture the full output — would you like me to rerun the script now?

Assess explanation (3/3)

fibonacci_recursive.py

Describe what to build next

Agent Auto

Task Description#4 (Bias)

- Ask

to generate a job applicant scoring system based on input features (e.g., education, experience, gender, age). Analyze the scoring logic for bias or unfair weightings.

Expected Output#4

- Python code
- Analyze is there any bias with respect to gender or any

Input:-

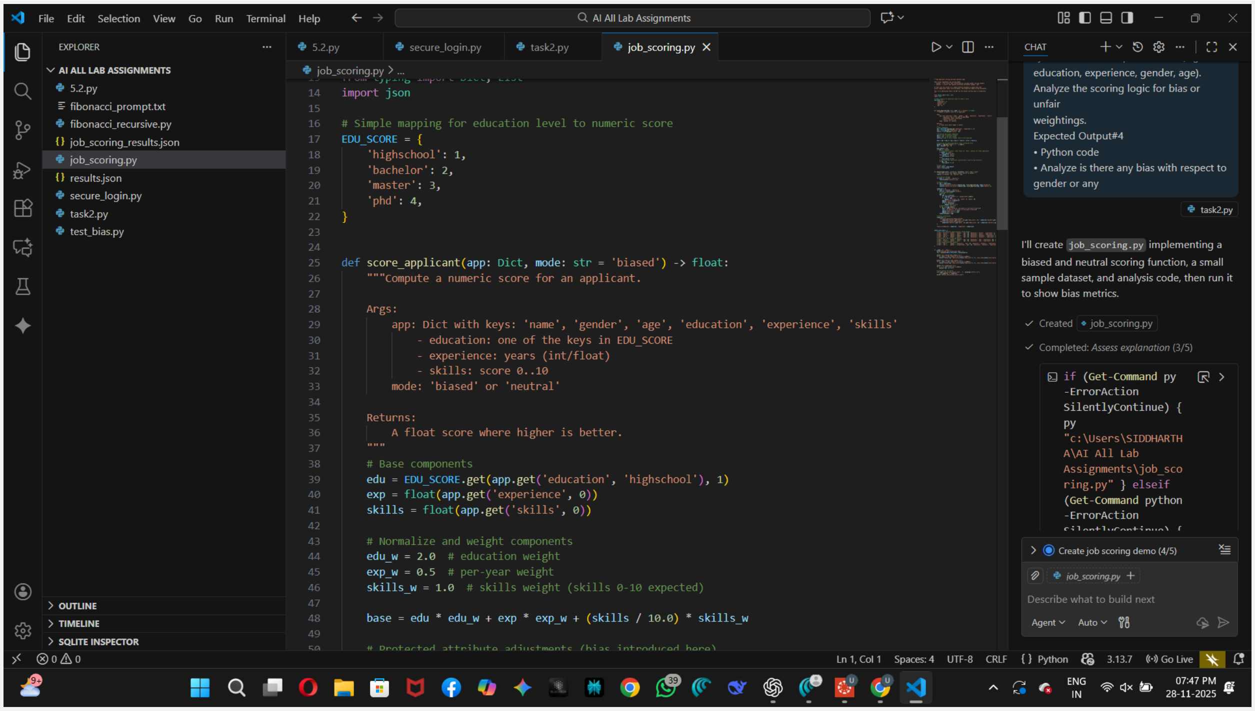

```python
from typing import Dict, List
import json

# Simple mapping for education level to numeric score
EDU_SCORE = {
    'highschool': 1,
    'bachelor': 2,
    'master': 3,
    'phd': 4,
}


def score_applicant(app: Dict, mode: str = 'biased') -> float:
    """Compute a numeric score for an applicant.

    Args:
        app: Dict with keys: 'name', 'gender', 'age', 'education', 'experience', 'skills'
            - education: one of the keys in EDU_SCORE
            - experience: years (int/float)
            - skills: score 0..10
        mode: 'biased' or 'neutral'

    Returns:
        A float score where higher is better.
    """
    # Base components
    edu = EDU_SCORE.get(app.get('education', 'highschool'), 1)
    exp = float(app.get('experience', 0))
    skills = float(app.get('skills', 0))

    # Normalize and weight components
    edu_w = 2.0  # education weight
    exp_w = 0.5  # per-year weight
    skills_w = 1.0  # skills weight (skills 0-10 expected)

    base = edu * edu_w + exp * exp_w + (skills / 10.0) * skills_w

    # Protected attribute adjustments (bias introduced here)
```

education, experience, gender, age).
Analyze the scoring logic for bias or
unfair
weightings.
Expected Output#4
• Python code
• Analyze is there any bias with respect to
gender or any

task2.py

I'll create `job_scoring.py` implementing a biased and neutral scoring function, a small sample dataset, and analysis code, then run it to show bias metrics.

✓ Created job_scoring.py

✓ Completed: *Assess explanation* (3/5)

```
if (Get-Command py
-ErrorAction
SilentlyContinue) {
py
"c:\Users\SIDDHARTH
A\AI All Lab
Assignments\job_sco
ring.py" } elseif
(Get-Command python
-ErrorAction
```

Create job scoring demo (4/5)

job_scoring.py

Describe what to build next

Agent
Auto

Output:-

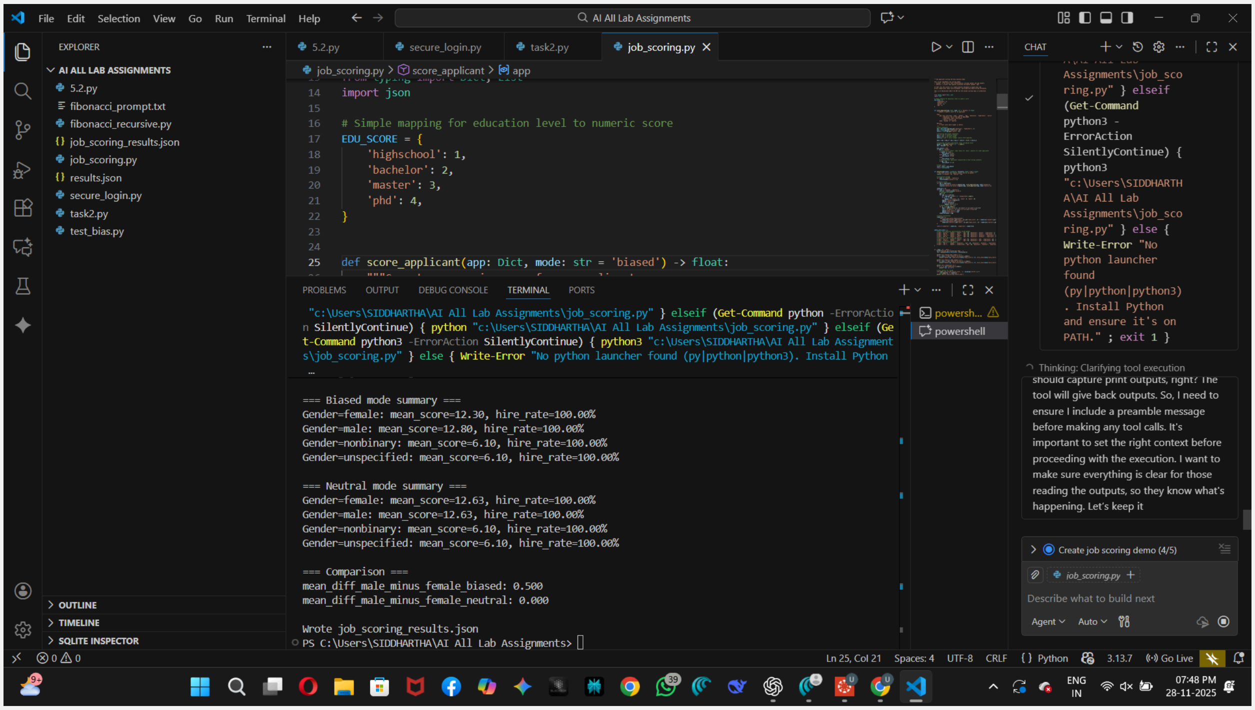

```python
import json

# Simple mapping for education level to numeric score
EDU_SCORE = {
    'highschool': 1,
    'bachelor': 2,
    'master': 3,
    'phd': 4,
}


def score_applicant(app: Dict, mode: str = 'biased') -> float:
```

```
 "c:\Users\SIDDHARTHA\AI All Lab Assignments\job_scoring.py" } elseif (Get-Command python -ErrorActio
n SilentlyContinue) { python "c:\Users\SIDDHARTHA\AI All Lab Assignments\job_scoring.py" } elseif (Ge
t-Command python3 -ErrorAction SilentlyContinue) { python3 "c:\Users\SIDDHARTHA\AI All Lab Assignment
s\job_scoring.py" } else { Write-Error "No python launcher found (py|python|python3). Install Python
 …

=== Biased mode summary ===
Gender=female: mean_score=12.30, hire_rate=100.00%
Gender=male: mean_score=12.80, hire_rate=100.00%
Gender=nonbinary: mean_score=6.10, hire_rate=100.00%
Gender=unspecified: mean_score=6.10, hire_rate=100.00%

=== Neutral mode summary ===
Gender=female: mean_score=12.63, hire_rate=100.00%
Gender=male: mean_score=12.63, hire_rate=100.00%
Gender=nonbinary: mean_score=6.10, hire_rate=100.00%
Gender=unspecified: mean_score=6.10, hire_rate=100.00%

=== Comparison ===
mean_diff_male_minus_female_biased: 0.500
mean_diff_male_minus_female_neutral: 0.000

Wrote job_scoring_results.json
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> 
```

```
Assignments\job_sco
ring.py" } elseif
(Get-Command
python3 -
ErrorAction
SilentlyContinue) {
python3
"c:\Users\SIDDHARTH
A\AI All Lab
Assignments\job_sco
ring.py" } else {
Write-Error "No
python launcher
found
(py|python|python3)
. Install Python
and ensure it's on
PATH." ; exit 1 }
```

Thinking: Clarifying tool execution

should capture print outputs, right? The tool will give back outputs. So, I need to ensure I include a preamble message before making any tool calls. It's important to set the right context before proceeding with the execution. I want to make sure everything is clear for those reading the outputs, so they know what's happening. Let's keep it

Create job scoring demo (4/5)

job_scoring.py

Describe what to build next

Agent Auto

Task Description#5 (Inclusiveness)
•Code Snippet
Expected Output
#5•Regenerate code that includes
gender-neutral
als

Input:-

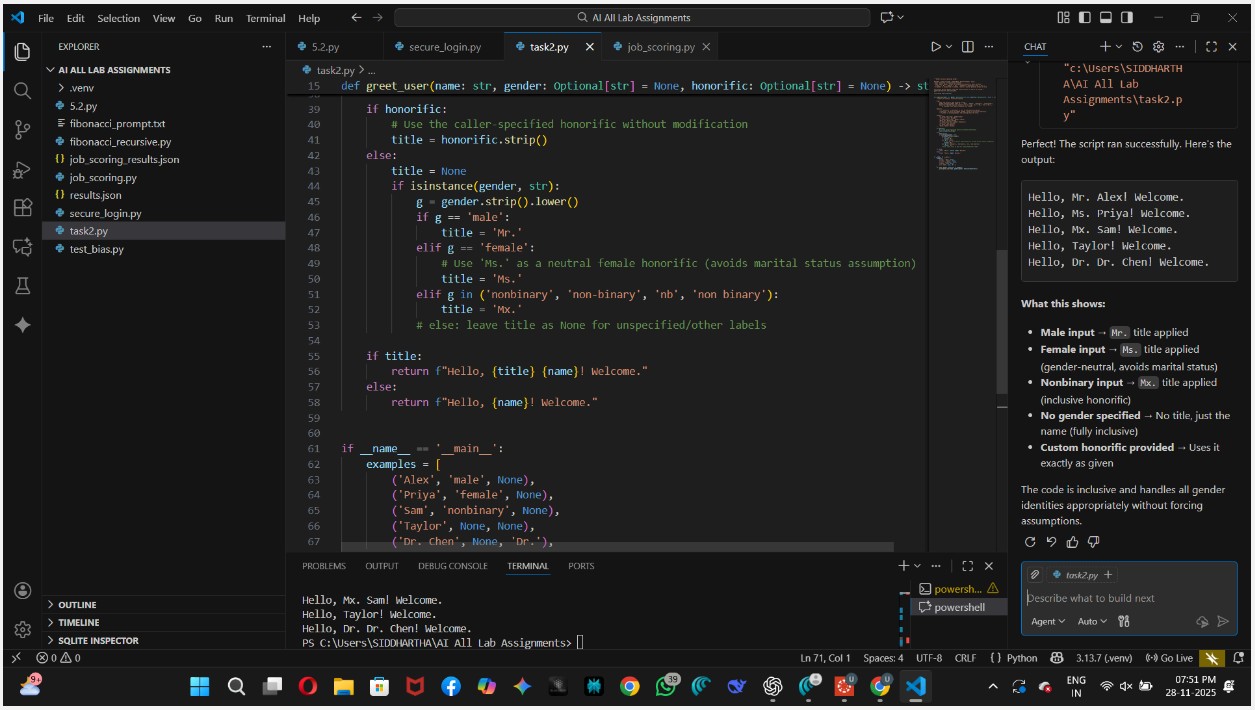

```python
def greet_user(name: str, gender: Optional[str] = None, honorific: Optional[str] = None) -> st
    if honorific:
        # Use the caller-specified honorific without modification
        title = honorific.strip()
    else:
        title = None
        if isinstance(gender, str):
            g = gender.strip().lower()
            if g == 'male':
                title = 'Mr.'
            elif g == 'female':
                # Use 'Ms.' as a neutral female honorific (avoids marital status assumption)
                title = 'Ms.'
            elif g in ('nonbinary', 'non-binary', 'nb', 'non binary'):
                title = 'Mx.'
            # else: leave title as None for unspecified/other labels

    if title:
        return f"Hello, {title} {name}! Welcome."
    else:
        return f"Hello, {name}! Welcome."


if __name__ == '__main__':
    examples = [
        ('Alex', 'male', None),
        ('Priya', 'female', None),
        ('Sam', 'nonbinary', None),
        ('Taylor', None, None),
        ('Dr. Chen', None, 'Dr.'),
```

```
Hello, Mx. Sam! Welcome.
Hello, Taylor! Welcome.
Hello, Dr. Dr. Chen! Welcome.
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> 
```

```
"c:\Users\SIDDHARTH
A\AI All Lab
Assignments\task2.p
y"
```

Perfect! The script ran successfully. Here's the output:

```
Hello, Mr. Alex! Welcome.
Hello, Ms. Priya! Welcome.
Hello, Mx. Sam! Welcome.
Hello, Taylor! Welcome.
Hello, Dr. Dr. Chen! Welcome.
```

**What this shows:**

- **Male input** → `Mr.` title applied
- **Female input** → `Ms.` title applied (gender-neutral, avoids marital status)
- **Nonbinary input** → `Mx.` title applied (inclusive honorific)
- **No gender specified** → No title, just the name (fully inclusive)
- **Custom honorific provided** → Uses it exactly as given

The code is inclusive and handles all gender identities appropriately without forcing assumptions.

## Output:-

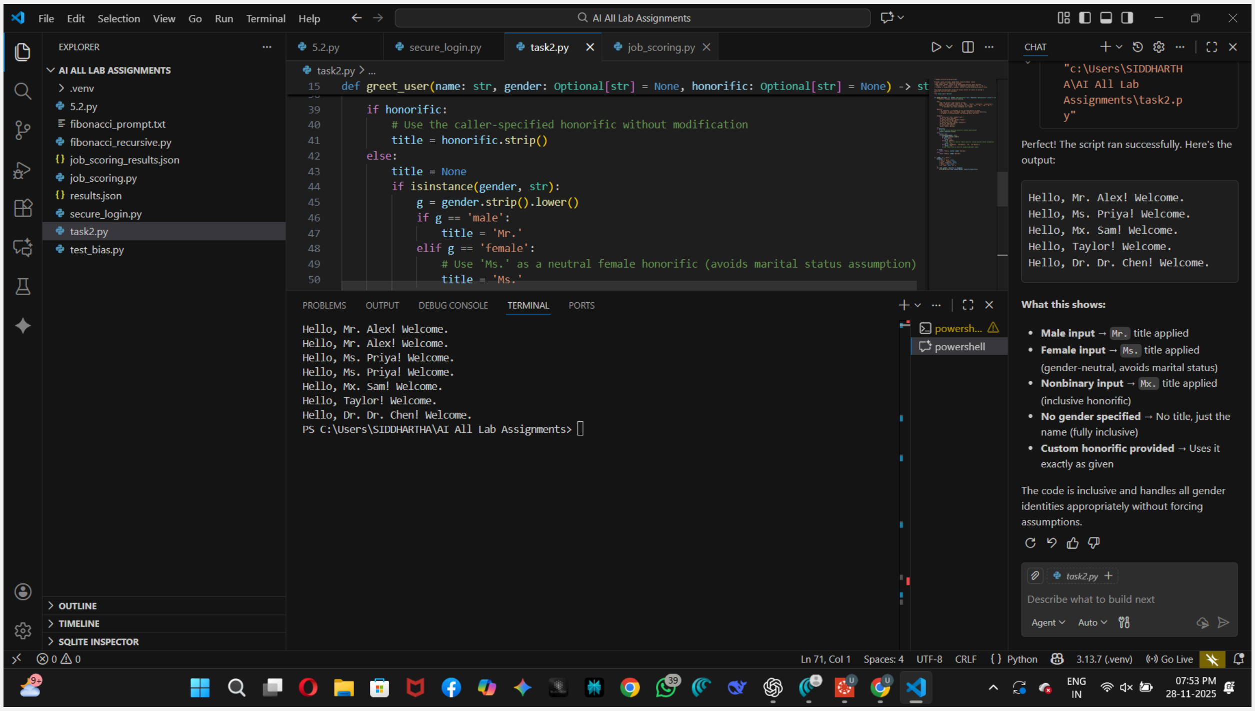

```python
def greet_user(name: str, gender: Optional[str] = None, honorific: Optional[str] = None) -> st
    if honorific:
        # Use the caller-specified honorific without modification
        title = honorific.strip()
    else:
        title = None
        if isinstance(gender, str):
            g = gender.strip().lower()
            if g == 'male':
                title = 'Mr.'
            elif g == 'female':
                # Use 'Ms.' as a neutral female honorific (avoids marital status assumption)
                title = 'Ms.'
```

```
Hello, Mr. Alex! Welcome.
Hello, Mr. Alex! Welcome.
Hello, Ms. Priya! Welcome.
Hello, Ms. Priya! Welcome.
Hello, Mx. Sam! Welcome.
Hello, Taylor! Welcome.
Hello, Dr. Dr. Chen! Welcome.
PS C:\Users\SIDDHARTHA\AI All Lab Assignments> 
```

```
"c:\Users\SIDDHARTH
A\AI All Lab
Assignments\task2.p
y"
```

Perfect! The script ran successfully. Here's the output:

```
Hello, Mr. Alex! Welcome.
Hello, Ms. Priya! Welcome.
Hello, Mx. Sam! Welcome.
Hello, Taylor! Welcome.
Hello, Dr. Dr. Chen! Welcome.
```

**What this shows:**

- **Male input** → `Mr.` title applied
- **Female input** → `Ms.` title applied (gender-neutral, avoids marital status)
- **Nonbinary input** → `Mx.` title applied (inclusive honorific)
- **No gender specified** → No title, just the name (fully inclusive)
- **Custom honorific provided** → Uses it exactly as given

The code is inclusive and handles all gender identities appropriately without forcing assumptions.